

Room Booking - High Level Design

- [Introduction](#)
 - [Overview](#)
 - [Scope](#)
 - [Out of Scope](#)
- [Requirements](#)
 - [User](#)
 - [Room Search](#)
 - [Booking](#)
 - [Api](#)
- [Components diagram](#)
- [5\) Database Schema \(MongoDB\)](#)
 - [Collections & Schema](#)
 - [users Collection](#)
 - [rooms Collection](#)
 - [availability Collection](#)
 - Option A — Date-based subdocument (recommended for fast queries):
 - Mongo TTL or write-heavy warning: For very large calendars (years), we avoid huge documents.
 - [bookings Collection](#)
 - Concurrency Handling with MongoDB
 - A. Atomic Single-Document Updates (Recommended for per-day availability)
 - B. Multi-document Transactions (for multi-day bookings)
 - [5. Scalability Strategies](#)

PAGE UNDER CONSTRUCTION

Introduction

Overview

This document describes the high level design of the room-booking platform that enables users to register (login) , search for available rooms, make or cancel bookings and retrieve their booking history in real time using API-driven microservices, multi-region databases, caching, and strong concurrency controls to prevent double-booking

Scope

The design aims to drill down the room-booking architecture topic and elaborate on the technical details required for implementation/integration.

Out of Scope

There is an external system responsible for inserting room data and availability into the database, as well as removing unavailable rooms.

In addition, a background job invalidates the cache based on updates from the real database.

Requirements

User

- Users can **register**, and **log in**,
- Users can **view their booking history**.

Room Search

- Search rooms by **date**, **location**, **capacity**, **price**, and **amenities**.
- Real-time **availability check**.

Booking

- Users can **create a booking** for a room.

- Users can **cancel a booking**.
- System must **prevent double-booking**.
- Users receive **confirmation notifications** (email/SMS)

Api

User:

The user API must implement a “forgot password” feature as well.
The validation of the email format and password must be handled in the UI (and enforce in BE we as well).

POST /api/v1/users/register

Headers

- Content-Type: application/json

body:

```
{
  "email": "yosefy2@hotmail.com",
  "password": "Lavi:99798e",
  "fullName": "Yaniv yosef",
  "phone": "+972500000000",
  "locale": "en-IL"
}
```

Responses

Status code

- 201 - successfully register created


```
{ "userId": "u_123456", "email": "yosefy2@hotmail.com", "createdAt": "2025-11-26T09:00:00Z", token:token }
```
- 400 - missing or invalid body
- 409 - conflict - user already exists
- 500 - internal server error

POST /api/v1/users/login

Headers

- Content-Type: application/json

body:

```
{
  "email": "yosefy2@hotmail.com",
  "password": "Lavi:99798e",
}
```

Responses

Status code

- 200 - successfully logIn


```
{ userId: "u_123456" , token:token }
```
- 400 - missing or invalid body
- 500 - internal server error .

GET /api/v1/users/{userId}/bookingsHistory
Query Parameters : startDate , endDate

Headers

- Content-Type: application/json
- Authorization: Bearer <access_token>

Responses

Status code

- 200 - Success
 { bookingsHistory:bookingsHistory }
- 404 - no user Id
- 500 - internal server error .

Booking:

POST /api/v1/bookings

Headers

- Content-Type: application/json
- Authorization: Bearer <access_token>

body:

```
{
  "id": "123",
  "roomId": "r_987",
  "userId": "u_123456",
  "checkin": "2025-12-20",
  "checkout": "2025-12-22",
  "guests": 2,
  "payment": {
    "method": "card",
    "paymentToken": "tok_visa_123"
  },
  "metadata": { "source": "web" }
}
```

Responses

Status code

- 201 Created (booking confirmed)
 { "roomId": }
- 400 - missing or invalid body
- 409 Conflict (double-booking conflict)
- 422 Unprocessable Entity (invalid dates)
- 429 Too Many Requests (rate-limited)
- 500 internal server error

Note: Display a warning before calling this API to indicate that the cancellation complies with policy
POST /api/v1/bookings/{bookingId}/cancel

Headers

- Content-Type: application/json
- Authorization: Bearer <access_token>

body:

```
{
  "reason": "sick",
}
```

Responses

Status code

- 201 Created (booking confirmed)
 { "roomId": }
- 403 Forbidden (user does not own this booking)
- 404 Booking not found
- 409 Cannot cancel (e.g., past booking, outside cancellation window)
- 500 Internal server error
- 401 Unauthorized (invalid/missing token)

search room:

GET /api/v1/rooms/search

Query Parameters : startDate , endDate <TBD>

Headers

- Content-Type: application/json
- Authorization: Bearer <access_token>

Responses


Status code

- 200 ok


```
{
  "rooms": [
    {
      ...
    }
  ]
}
```
- 400 - invalid body Query Parameters
- 429 Too Many Requests (rate-limited)
- 500 internal server error

Components diagram

The below diagram demonstrates the Room-booking BE.



5) Database Schema (MongoDB)

Document-oriented, denormalized where useful, with careful use of references for transactional parts (e.g., bookings). Includes indexing strategy and consistency considerations.

Collections & Schema

users Collection

Purpose: Authentication, profile, preferences.

Example document:

```
{ "_id": ObjectId("..."), "email": "user@example.com", "passwordHash": "...", "name": "John Doe", "phone": "+1-555-1234", "createdAt": ISODate("2025-01-01T10:00:00Z"), "updatedAt": ISODate("2025-01-05T08:00:00Z") }
```

Indexes:

- { email: 1 } — **unique**

- { createdAt: -1 } — listing & admin analytics
-

rooms Collection

Purpose: Static room info (hotel, type, capacity, location).
Stored separately from availability to reduce document growth.

```
{ "_id": ObjectId(...), "hotelId": ObjectId(...), "type": "Deluxe Suite", "capacity": 4, "price": 180, "location": { "type": "Point", "coordinates": [34.78, 32.07] }, "images": ["s3://bucket/1.jpg", "s3://bucket/2.jpg"], "amenities": ["wifi", "parking", "pool"], "createdAt": ISODate(), "updatedAt": ISODate() }
```

Indexes:

- { location: "2dsphere" } — geo-search
- { hotelId: 1 }
- { price: 1, capacity: 1 }

availability Collection

Purpose: Fast lookup of available dates per room.

Two models possible choose per load requirements:

Option A — Date-based subdocument (recommended for fast queries):

```
{ "_id": ObjectId(...), "roomId": ObjectId(...), "dates": { "2025-06-20": "available", "2025-06-21": "reserved", "2025-06-22": "blocked" }, "updatedAt": ISODate() }
```

Mongo TTL or write-heavy warning: For very large calendars (years), we avoid huge documents.

Indexes:

- { roomId: 1 }
 - Or if using "1 doc per date": { roomId: 1, date: 1 }
-

bookings Collection

Purpose: Main transactional entity.
A booking is immutable after confirmation (eventually written to Kafka search index).

```
{ "_id": ObjectId(...), "userId": ObjectId(...), "roomId": ObjectId(...), "hotelId": ObjectId(...), "startDate": "2025-06-21", "endDate": "2025-06-24", "status": "confirmed", // pending | confirmed | canceled "paymentId": ObjectId(...), "createdAt": ISODate(), "updatedAt": ISODate() }
```

Indexes:

- { userId: 1, createdAt: -1 } — history view
- { roomId: 1, startDate: 1, endDate: 1 } — conflict detection
- { status: 1 } — operational filters / queueing

Concurrency Handling with MongoDB

A. Atomic Single-Document Updates (Recommended for per-day availability)

Model:

- Collection: room_calendar
- One document per (room_id, date) with fields:

```
{ "_id": ObjectId(), "room_id": "room123", "date": "2025-12-01", "is_available": true, "inventory": 5 }
```

Booking flow (atomic update):

```

const bookingResult = await db.collection('room_calendar').findOneAndUpdate( { room_id: roomId, date: bookingDate, is_available: true }, { $set: { is_available: false } }, // or decrement inventory: { $inc: { inventory: -1 } } { returnDocument: 'after' } ); if (!bookingResult.value) { throw new Error('Room not available') }
;

```

Guarantees:

- Atomic per document.
- Avoids double-booking for single-day per-room entries.

B. Multi-document Transactions (for multi-day bookings)

Model:

- Collection: `room_calendar` with one document per `(room_id, date)`.

Booking flow:

```

const session = client.startSession(); try { await session.withTransaction(async () => { const docs = await db.collection('room_calendar').find({ room_id: roomId, date: { $gte: startDate, $lte: endDate } }).toArray(); if (docs.some(doc => !doc.is_available)) { throw new Error('Some dates not available'); } await db.collection('room_calendar').updateMany( { room_id: roomId, date: { $gte: startDate, $lte: endDate } }, { $set: { is_available: false } }, { session } ); await db.collection('bookings').insertOne({ room_id: roomId, start_date: startDate, end_date: endDate, user_id: userId, created_at: new Date() }, { session }); } finally { await session.endSession(); }

```

Guarantees:

- Strong consistency for multi-day bookings in replica sets or sharded clusters.
- Avoids double-booking across multiple documents.

5. Scalability Strategies

For Search

- Use **Elasticsearch / OpenSearch** cluster as before for full-text search, indexing room documents and precomputing availability snapshots.
- Sync from MongoDB via **change streams** for near-real-time updates instead of Debezium/Kafka.
- Use caches and CDN for static lists (popular queries).
- Use **geo-sharding** or **hashed sharding** in Elasticsearch indices if needed.

For Booking throughput

- Keep the **booking service stateless** and scale horizontally.
- For **write consistency**, MongoDB supports **multi-document transactions** in replica sets.
 - Option 1 (simpler): All writes go to a single primary (leader) in the primary region.
 - Option 2 (advanced): For multi-region setup, use **MongoDB global clusters** with **zone sharding** to route writes to the correct region while maintaining strong consistency via transactions.
- Use **MongoDB change streams** or a **message bus (Kafka, RabbitMQ)** for async tasks like notifications or analytics.

Caching

- Use **Redis** for hot-room details and per-room availability caches.
- Use short TTLs and proactive invalidation on booking events.
- Consider caching query results for frequent availability checks.

Multi-region

- Deploy services to multiple regions for low read latency and high availability.
- Option 1: Single-primary writes; reads from secondary replicas. Simple, consistent.
- Option 2: MongoDB global clusters with **zone sharding** for multi-region writes and strong consistency.

Autoscaling & resilience

- Use **HPA** for services; implement circuit breakers and health checks.
- Retry with jitter for transient errors.
- Bulkhead isolation: separate booking service from analytics to prevent noisy neighbors.
- Use feature flags and gradual rollout gates.

Monitoring & SLOs

- Track request latency percentiles, error rates, booking conflict rate, queue depth.
- Maintain **SLOs** with error budgets.
- Monitor MongoDB metrics: replication lag, transaction conflicts, operation throughput.

